Overnight run — 6 September 2026

Generated Saturday 05 September, 23:52. Every number below is read from the files on disk at render time, so a phase that did not run cannot appear here as though it did.

What ran

track	phase	outcome	time
gemma31	A-calibration	**running**	—
qwen	A-calibration	**running**	—
gemma26	A-calibration	**running**	—

The main question: which graders can see near the ceiling?

A grader with no spread reports a difference of zero whatever is true. Yesterday that nearly published four arm-order reversals that did not exist. This measures it directly: the same 60 answers, spanning the score range, graded by every model.

model	ceiling band	mid band	low band	usable near ceiling?
gemma31	—	—	—	*did not run*
qwen	—	—	—	*did not run*
gemma26	—	—	—	*did not run*

"distinct" is how many different scores the model gave across 20 answers. A model that gave the same score to all 20 cannot tell them apart, and any comparison built on it is empty.

Does a third grader agree that writing your own exam flatters you?

Not measured. did not run

That is a result rather than a gap: if no available model can tell these answers apart, the claim rests on one grader for a reason we can state, and it belongs in the paper's limitations.

What this night does not settle

• Nothing here re-measures the system. No questions were re-answered except the 30-question probe; every other number is a re-grading of answers that already existed.

• A grader that agrees is not proof of correctness. All of this tests whether conclusions survive a change of judge, not whether the judge is right.

• Two known scoring faults remain deliberately unfixed, because fixing them mid-comparison would break the comparison. Both would raise scores.