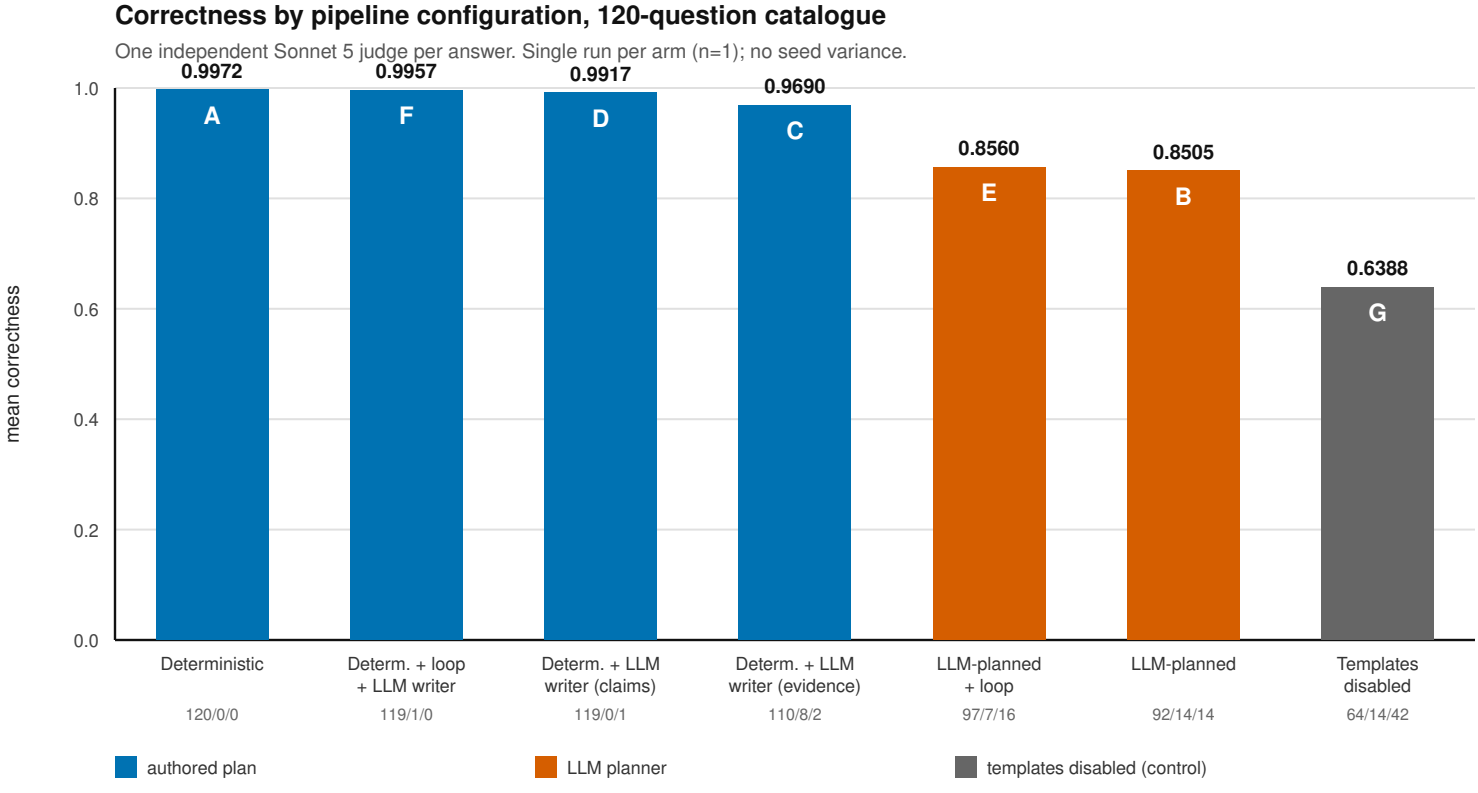
Correctness by pipeline configuration, 120-question catalogue
One independent Sonnet 5 judge per answer. Single run per arm (n=1); no seed variance.
mean correctness
1.0
0.8
0.6
0.4
0.2
0.0
0.9972
A
0.9957
F
0.9917
D
0.9690
C
0.8560
E
0.8505
B
0.6388
G
Deterministic
120/0/0
Determ. + loop + LLM writer
119/1/0
Determ. + LLM writer (claims)
119/0/1
Determ. + LLM writer (evidence)
110/8/2
LLM-planned + loop
97/7/16
LLM-planned
92/14/14
Templates disabled
64/14/42
authored plan
LLM planner
templates disabled (control)
Bar labels: mean correctness. Under each bar: CORRECT/PARTIAL/INCORRECT. Generated by figures/sources/make_arm_comparison.py