Overnight report — 2026-09-07

Driver started **2026-09-07 00:49:19.826761**, deadline **2026-09-07 08:00:00**.

Endpoint: qwen38-27B-v2 on :9090, sharded across **default-model-2 / -3 / -4** at 32 workers each. **default-model was reserved by the operator at 01:15 and is untouched from that point.**

> **What did not happen.** The driver exited at 02:49 having completed every step it was given, not at the 08:00 deadline. Its steps finished early because judging completed early. No replacement was scheduled, so between 04:00 and 09:57 nothing ran: three extension seeds generated at ~03:30 sat ungraded, and no 08:00 report was produced. Grading started 09:57. The gap is an operator error, not a tool failure -- there was a watchdog for judging and none for the schedule.

Steps

step	outcome	elapsed
judge_generic_resume_1	rc=143	17 min
judge_generic_resume_2	rc=143	1 min
judge_generic_resume_3	ok	95 min
validate_research	ok	0 min
verify_paper_numbers	ok	0 min
identity_analysis	ok	0 min
report	ok	0 min

EXP-2026-09-011 — does generic's wording variation move its score?

This is the cell the night was for. EXP-2026-09-010 established that generic words 143 of 276 answers differently on a rerun at the same seed, in the field the grader reads. Whether that reaches the **score** is this measurement.

seed	block q1	n	block q2	n
1729	0.6646	284	0.6394	284
20260906	0.6451	284	0.6444	284
424242	0.6468	284	0.6394	284

Extension seeds (EXP-2026-09-011 A-1)

seed	answers generated	judgments
31337	284	0
8675309	284	0
271828	284	0

Generated but ungraded. Three further seeds were generated to turn a three-point spread into something closer to a distribution, and grading was not queued behind the generation. The answers are frozen on disk and can be graded at any time; nothing is lost but the elapsed hours.

Balanced across-seed spread: 0.0039 (means [0.652, 0.6447, 0.6431]).

• **H2** (SD < 0.03): **holds**

• **H3, as pre-registered** (holdout gap 0.1684 > 3xSD = 0.0116): **holds**

• **H3, same population** (dev gap 0.1623 > 3xSD = 0.0116): **holds**

The pre-registered form of H3 compares a holdout gap against a spread measured on the dev half. That mismatch is in the protocol as written, so it is reported as written *and* alongside the same-population comparison, which is the one that actually tests what H3 means.

Counterbalancing note: both blocks are the **same grader**, so this is position-balanced repeatability, not cross-grader agreement. The second grader family was unavailable (its slot was freed).

Record

44 of 45 experiments complete.

Not complete:

• EXP-2026-07-010-independent-eval-baseline — superseded

Paper number traceability

• traced to a manifest: **86**

• untraced: **0**

This check did not exist before tonight. OUTLINE.md and the draft's own appendix both claimed automated enforcement of number→experiment traceability; a readiness audit found no such check. It is a **review list, not a gate**.

Generated 2026-09-07 09:58. Every figure is read from disk at render time; nothing here is carried over from prose.